

Surface mass balance inference based on a transient glacier evolution model inversion

Kejdi Lleshi¹, G. Jouvét¹, F. Herman¹, M. Kneib^{2,3}, and T. Gregov¹

¹Institute of Earth Surface Dynamics, University of Lausanne, Lausanne, Switzerland

²Laboratory of Hydraulics, Hydrology and Glaciology (VAW), ETH Zurich, Zurich, Switzerland

³Swiss Federal Institute for Forest, Snow and Landscape Research (WSL), Sion, Switzerland

Correspondence: Kejdi Lleshi (kejdi.lleshi@unil.ch)

Abstract. The distributed surface mass balance (SMB) of mountain glaciers is the key forcing for glacier evolution and for projections of meltwater availability. Yet direct observations remain confined to sparse stake networks, and distributed fields have so far been derived only by recent continuity approaches, which combine satellite-derived elevation change with surface velocity. These approaches necessarily assume a stationary ice flow over decadal timescales, and thereby ignore the slow-down of glaciers in a warming climate. Here, we present a new SMB inference method, implemented within the open-source Instructed Glacier Model (IGM), that recasts the problem as a fully differentiable transient inversion. The scheme uses observed surface velocities to initialise a dynamically consistent glacier model, then inverts multi-temporal surface elevation changes to estimate the spatial distribution of mass balance that best explains the observed glacier evolution. We apply the method to three well-studied glaciers in the Alps—Argentière, Great Aletsch and Rhône—evaluating all three with a single, glacier-independent pipeline. Validation against in situ stake measurements shows that surface velocity alone—with no ice-thickness data—suffices to recover a good fit to the observed mass balance. This suggests that the surface-velocity misfit offers a selection criterion where stakes are absent. These results validate our method and render it transferable to any glacier with accurate multi-temporal DEMs and surface velocities, even in the absence of stake measurements, opening a path towards regional-scale, distributed SMB inference.

1 Introduction

Mountain glaciers are sensitive recorders of climate and a critical component of regional water resources, yet projecting their evolution requires computing their distributed surface mass balance (SMB)—the spatial pattern of accumulation and ablation—which remains poorly constrained at the global scale, particularly in remote parts of the world. This is in part because the direct measurement of SMB relies on stake networks and snow pits that are labour-intensive and hazardous in steep accumulation areas. They are consequently deployed on less than 0.2 % of the world’s glacierised area (Cogley et al., 2011). The resulting records, compiled by the World Glacier Monitoring Service (WGMS, 2025), cover only a few accessible benchmark glaciers (Zemp et al., 2019). Large-scale glacier models therefore rarely calibrate on stake data, which serve mainly as an independent check; instead, the parameters of their SMB models are tuned against glacier-wide geodetic mass balances (Marzeion et al., 2012; Radić et al., 2014; Hugonnet et al., 2021). As remote-sensing products of glacier surface velocity (Millan et al., 2022)

25 and elevation change (Hugonnet et al., 2021; Beraud et al., 2022) reach near-global coverage, the question is no longer whether the data exist, but how to combine them to infer physically consistent SMB fields.

1.1 Existing approaches

SMB is most commonly obtained from glacier models that quantify the accumulation and ablation terms from meteorological inputs—at least precipitation and temperature, and further variables where a surface energy balance is resolved (Marzeion et al., 2012; Rounce et al., 2020a; Buri et al., 2023). These inputs come from climate models and reanalyses whose resolution is too coarse to represent high-mountain precipitation and temperature accurately, so the forcing must be bias-corrected. With calibration limited to glacier-wide geodetic balances, the problem stays strongly underdetermined: different parameter sets and model formulations reproduce the same integrated balance while implying different distributed SMB patterns and diverging projections (Rounce et al., 2020b; Schuster et al., 2023).

35 Assimilation frameworks tackle this by calibrating against distributed observations rather than a glacier-wide total: FROST (Herrmann et al., 2025) tunes an elevation-dependent SMB model coupled to IGM ice dynamics to infer equilibrium-line altitudes across the Alps, fitting the parameters of a prescribed SMB form with an ensemble Kalman filter. Being derivative-free, it obtains parameter sensitivities from the spread of an ensemble that must grow with the number of unknowns, whereas differentiating end to end through our forward model returns the sensitivity to every parameter at roughly the cost of one additional forward run. Alternatively, sparse direct measurements are extended to distributed fields by homogenising them against decadal geodetic volume change, as in the long-running Glacier Monitoring Switzerland programme (GLAMOS; GLAMOS, 2021). These provide a valuable observational reference, but rest on a large number of statistical assumptions and, more importantly, exist for only a handful of glaciers.

Satellite-derived velocity and elevation-change fields are now widely available. This has driven a family of approaches, closest to the present work, that infer the SMB from mass-conservation principles. The governing equation relates the rate of surface-elevation change to the divergence of the ice flux and the SMB (Cuffey and Paterson, 2010): given ice thickness and surface velocity, the vertically integrated flux can be computed, differentiated spatially, and combined with an observed elevation-change field to provide an SMB estimate. From early flux-gate estimates of emergence velocity in the ablation zone (Berthier and Vincent, 2012; Brun et al., 2018; Kneib et al., 2022), the approach has been generalised to fully distributed SMB over entire glaciers (Bisset et al., 2020; Miles et al., 2021; Pelto and Menounos, 2021; Van Tricht et al., 2021), to seasonal elevation-change time series (Zeller et al., 2023) and to debris-covered glaciers (Bhushan et al., 2024). A recent realisation is that of Kneib et al. (2024), who reconstructed a high-resolution distributed SMB for Argentière Glacier (French Alps) over 2012–2021 from Pléiades elevation change and velocity, ground-penetrating-radar thickness and three different ice-thickness models.

55 Two difficulties beset this route. First, the flux divergence is the spatial derivative of the product of two independently measured noisy fields—surface velocity and ice thickness—and is therefore dominated by small-scale noise, which is suppressed by *post hoc* spatial filtering (Van Tricht et al., 2021; Bisset et al., 2020). Filtering also serves a physical purpose, since averaging over a few ice thicknesses represents the longitudinal stress coupling that governs how ice flow responds to variations

in thickness and surface slope (Kamb and Echelmeyer, 1986; Enderlin et al., 2016); yet smoothing a quantity that respects
60 mass conservation locally redistributes mass between grid cells and violates the very balance the method enforces. Per-glacier
tuning of smoothing length scales, density assumptions and weighting schemes further hampers transfer to the regional scales
most relevant for water-resource applications. Second, the flux divergence is computed at a single observational snapshot and
assumed representative of the whole observation period, whereas glacier geometry, velocity and flux all evolve under ongoing
climate forcing (Jóhannesson et al., 1989; Zekollari and Huybrechts, 2015); treating the flux as stationary introduces a bias
65 that grows with the length of the window—a limitation Kneib et al. (2024) acknowledge explicitly. Transient assimilation
frameworks confront this directly, notably the dynamical spin-up employed by Hartl et al. (2025) and the AGILE framework
of Schmitt et al. (2026), which fit an evolving glacier model to observations rather than to a single snapshot.

1.2 This study

In this study we infer the SMB by inverting a transient glacier-evolution model: starting from a dynamically calibrated state,
70 the mass-conservation equation is integrated forward from one DEM to the next, and reverse-mode automatic differentiation
through the whole simulation returns the elevation-dependent SMB profile whose modelled geometry change best reproduces
the observed one. Inversions of transient ice-flow models are not themselves new—Brinkerhoff et al. (2025) jointly infer bed
elevation, basal traction and surface mass balance for Sit’ Tlein, with Bayesian quantification of the resulting uncertainty—but
they draw on rich, glacier-specific observations; our aim here is to recover the SMB from two DEMs and a single velocity field
75 alone.

Two features distinguish our formulation from flux-divergence inversions. First, the SMB is recovered by solving the forward
mass-conservation equation rather than by differentiating observed fields: the flux divergence is recomputed at every time step
from the evolving glacier state, so the inferred SMB is dynamically consistent over the whole window, rather than assuming
a stationary flux, and mass is conserved at every step, with no *post hoc* smoothing of a noisy divergence field. Second, we
80 parametrise the SMB as a one-dimensional function of surface elevation, evaluated at every grid cell to yield a 2D field;
elevation carries the first-order control on alpine SMB, so this leaves roughly ten unknowns in place of one per pixel and
constrains an inversion that would otherwise be underdetermined, with no length scale or density assumption left to fix per
site (Sect. 2.4). The price is that horizontal variability within an elevation band, such as the avalanche-fed hotspots resolved by
Kneib et al. (2024), is not recovered: a deliberate trade-off of sub-elevation detail for a transferable, tuning-free inversion, to
85 which we return in Sect. 4.

The accuracy of our geometry-based inversion is ultimately controlled by that of the input DEMs. Near-global elevation-
change products (Hugonnet et al., 2021) make the method broadly applicable, but their inaccuracies limit the fidelity of the
recovered SMB, a penalty we quantify directly by repeating one of the inversions on such a product (Appendix A1). Dedicated
multi-temporal DEMs deliver the highest accuracy, and the rapidly expanding archive of such data is what makes a regional-
90 scale application possible. We validate the method with two-epoch DEMs on three well-monitored Alpine glaciers spanning
 $\approx 12\text{--}80\text{ km}^2$ in area—Argentière (French Alps), Great Aletsch and Rhône (Swiss Alps)—all inferred by the *same* two-step

pipeline, without per-glacier retuning: the basal-motion parameter is reduced to a single scalar fixed by an ensemble search on the surface-velocity misfit, and no thickness survey is assumed.

Sect. 2 details the two-step pipeline—the input and validation data (Sect. 2.2), data assimilation of the initial dynamical state, then SMB inference by automatic differentiation through the transient forward run; Sect. 3 presents the inferred SMB profiles and their validation against in situ stake measurements; and Sect. 4 discusses the implications for regional-scale SMB inversion.

2 Methods

2.1 Overview

The observational inputs to our method are two surface digital elevation models (DEMs) of the glacier, acquired at epochs t_1 and t_2 several years apart, together with a surface-velocity field observed near t_2 , the epoch closer to the velocity acquisition for all three glaciers studied here. From these we infer the glacier’s surface mass balance (smb) with a two-step pipeline (Fig. 1), implemented within the open-source Instructed Glacier Model (IGM; Juvet et al., 2026). The elevation change between the two DEMs is the signal used to recover the smb (Sect. 2.2). At its core is the depth-integrated mass-conservation equation, which links this geometry change to the smb ,

$$\frac{\partial H}{\partial t} = smb - \nabla \cdot \mathbf{q}, \quad \mathbf{q} = H \bar{\mathbf{u}}, \quad (1)$$

where H is the ice thickness, smb is the surface mass balance, $\mathbf{q}$ is the flux and $\bar{\mathbf{u}}$ is the depth-averaged horizontal velocity. All mass terms are expressed in ice equivalent, so that (1) is dimensionally homogeneous; where the inferred smb is compared with stake measurements it is converted to water equivalent using an ice density of 917 kg m^{-3} . The velocity is supplied by IGM’s deep-learning ice-flow emulator (Juvet and Cordonnier, 2023; Rosier and Juvet, 2026), trained on a large ensemble of higher-order flow solutions to return $\bar{\mathbf{u}}$ from the ice geometry and the flow parameters. The emulator can be used in two ways. In the *online* mode its weights are updated during the simulation itself, at every step, by minimising the physical energy associated with the current geometry, so the flow model adapts continuously to the glacier at hand; because the weights then depend on the simulation history, the map from geometry to velocity is not a fixed one, which makes this mode ill-suited to inversion. In the *offline* mode the weights are pretrained on a controlled, high-quality catalogue of glacier geometries and then held frozen, so that map is a single deterministic function the inversion can differentiate through. We use the offline mode throughout, and refer the reader to Juvet et al. (2026) and Rosier and Juvet (2026) for details.

Equation (1) makes the central challenge explicit: the geometry evolution depends jointly on the unknown smb and on the ice dynamics through the flux $\mathbf{q}$. Inferring the smb from an observed geometry change therefore requires the dynamical state to be fixed beforehand; otherwise an error in the dynamics would be aliased into the inferred smb . An inconsistent initial state triggers an initialisation shock (Seroussi et al., 2011) that our short (up to ten-year) windows cannot absorb, unlike long spin-up simulations. Our pipeline reflects this dependence directly (Fig. 1): (i) we first calibrate the dynamical state by data assimilation against a surface-velocity snapshot near t_2 , inverting the ice-thickness field H and, with it, the bedrock (Sect. 2.3). (ii) Holding




Figure 1. Overview of the two-step SMB-inference workflow. *Step 1 (data assimilation)*: the observed surface velocity and the end-epoch DEM are assimilated to recover the ice thickness H_{t_2} —and hence the bedrock—while the scalar flow parameters (τ_{ref}, A) are selected by an ensemble search that minimises the surface-velocity misfit. *Step 2 (SMB inference)*: with the bedrock fixed, the start-epoch thickness H_{t_1} is derived from the observed start DEM, and the transient model is integrated from the start DEM to the end DEM under a candidate elevation-dependent profile $smb(z)$; reverse-mode automatic differentiation (AD) through the whole forward run yields the profile whose modelled end-of-window geometry best reproduces the calibrated one. The ice-flow emulator (frozen weights) supplies the depth-averaged velocity at every step.

the bedrock fixed, we derive the ice thickness at the start epoch from the observed t_1 DEM, then infer the smb : we integrate
 125 (1) forward from t_1 to t_2 and seek the smb whose simulated geometry best matches the calibrated one at t_2 (Sect. 2.4).

2.2 Data

We evaluate the method on three Alpine glaciers that span a wide range of size and are among the most intensively surveyed in the world (Table 1): Great Aletsch (Swiss Alps), the largest glacier of the Alps ($\approx 80 \text{ km}^2$); Rhône (Swiss Alps), a mid-sized valley glacier ($\approx 15 \text{ km}^2$) that feeds the eponymous river; and Argentière (French Alps, Mont-Blanc massif), a steep
 130 and fast-flowing valley glacier ($\approx 12 \text{ km}^2$). Each combines repeat DEMs, satellite surface-velocity fields and a decades-long in situ stake network—the Swiss elevation models from swisstopo (swisstopo, 2026), the stake records from GLAMOS in Switzerland (GLAMOS, 2023) and GLACIOCLIM in France (GLACIOCLIM, 2026)—which supply both the fields that drive the inversion and the ground truth that validates it.

Table 1. Input and validation data for the three study glaciers. DEM epochs bracket the inversion window; velocities, acquired near the end epoch t_2 for all three glaciers, drive the data assimilation there (Sect. 2.3). The stake column gives the number of continuous-record locations averaged for the stake RMSE.

	Argentière	Great Aletsch	Rhône
Time window	2012–2021	2009–2017	2007–2016
Surface DEMs	Pléiades	Swisstopo	Swisstopo
Velocity	(Millan et al., 2022)		
Stake locations	25	2	5

For each glacier the pipeline requires two surface DEMs bracketing the inversion window—whose difference is the geometry change to be explained—and a surface-velocity field for the data assimilation (Sect. 2.3). For Argentière the two DEMs are the 2012 and 2021 Pléiades elevation models of Kneib et al. (2024); for Great Aletsch (2009, 2017) and Rhône (2007, 2016) they are swisstopo geodetic DEMs. We resample all fields onto a common 100 m grid. Surface velocities are the satellite feature-tracking mosaics of Millan et al. (2022), built from Sentinel-2 and Landsat-8 image pairs over 2017–2018. For all three glaciers this snapshot lies several years closer to the end epoch t_2 than to t_1 (Table 1); we therefore anchor the data assimilation at t_2 (Sect. 2.3), making the thickness inversion contemporaneous with its own velocity observations rather than bridging the larger gap to t_1 . Alpine surface velocities also vary slowly relative to the decadal DEM windows (Dehecq et al., 2019), so the residual one- to four-year gap between the velocity snapshot and t_2 is expected to be immaterial.

For validation we use in situ stake balances, kept strictly independent of the inversion. Rather than compare against every raw stake reading—noisy and irregularly sampled in space and time—we retain only locations with a continuous multi-year record, average each into a single balance value ($\pm$ interannual standard deviation), and evaluate the inferred smb against these averages. This yields two such locations on Aletsch and five on Rhône; on Argentière the denser GLACIOCLIM network provides the published 2012–2021 mean balance at each of 25 stakes.

2.3 Ice-dynamics calibration via data assimilation

The first step of the pipeline uses two of the observational fields introduced above—the surface velocity and the t_2 DEM—to recover the ice-thickness field that places the glacier in a dynamically consistent state, free of the initialisation shock described in Sect. 2.1.

We use IGM’s existing data-assimilation module (Jouvet, 2023) to invert for the thickness H by minimising the regularised misfit

$$\mathcal{J}_{\text{DA}}(H) = \alpha_u \left\| \mathbf{u}_s(H) - \mathbf{u}_s^{\text{obs}} \right\|^2 + \mathcal{R}(H), \quad (2)$$

where $\mathbf{u}_s(H)$ is the modelled surface velocity returned by the emulator (Rosier and Jouvet, 2026), $\mathbf{u}_s^{\text{obs}}$ is the observed surface velocity (Millan et al., 2022), and $\mathcal{R}$ is a Tikhonov smoothness regulariser on H . The smoothness weight is set at the corner of




Figure 2. Selecting the smoothness weight of the data assimilation, shown for Great Aletsch. Each point is one assimilation, labelled by its Tikhonov weight: stronger smoothing lowers the thickness roughness (horizontal axis) but worsens the fit to the observed surface velocity (vertical axis). The corner of the curve (star) is the point beyond which further smoothing buys little reduction in roughness for a rapidly rising misfit, and fixes the weight used thereafter. The same sequence of weights and the same corner-detection rule are applied to all three glaciers, returning 100 (Argentière), 600 (Great Aletsch) and 1000 (Rhône).

the resulting L-curve (Hansen, 1992), beyond which further smoothing degrades the velocity fit without appreciably reducing thickness roughness (Fig. 2). The minimisation is solved by L-BFGS (Liu and Nocedal, 1989); we refer the reader to Jouvét (2023) for the full formulation, weighting strategy and convergence behaviour.

160 The emulated flow regime is not set by the geometry alone: it also depends on two scalar physical parameters, the basal-motion parameter τ_{ref} and the Arrhenius factor A . Here τ_{ref} (in MPa) is the reference basal shear stress in Weertman’s sliding law (Weertman, 1957)—the shear stress sustaining a reference basal speed of 100 m yr^{-1} —so that a smaller τ_{ref} corresponds to easier sliding. Rather than invert these pointwise, we fix them beforehand by an ensemble search: candidate (τ_{ref}, A) pairs are drawn log-uniformly from $[0.05, 2.0] \text{ MPa}$ and $[30, 200] \text{ MPa}^{-3} \text{ yr}^{-1}$ by a Tree-structured Parzen Estimator (TPE; Bergstra et al., 2011), as implemented in Optuna (Akiba et al., 2019), and each pair is passed through the data assimilation; we retain the pair that minimises the surface-velocity misfit, so the assimilation itself inverts the thickness H alone. We apply this identical procedure to all three glaciers, which keeps the pipeline free of glacier-specific tuning.

Data assimilation inverts the thickness at the epoch closest to the velocity observations—the end epoch t_2 for all three glaciers studied here—which we denote H_{t_2} . Finding the thickness fixes the bedrock,

$$170 \quad z_b(\mathbf{x}) = z_s^{\text{obs}}(\mathbf{x}, t_2) - H_{t_2}(\mathbf{x}), \quad (3)$$

which provides the substrate on which the transient run subsequently evolves: held fixed, it also gives the thickness at the start epoch t_1 directly from the observed t_1 DEM (Sect. 2.4). Only once the calibrated state matches the observed velocity to within observational uncertainty do we accept it and proceed to the SMB inversion.

2.4 SMB inference by automatic differentiation

175 2.4.1 Initial condition and target.

With the dynamical state calibrated at t_2 (Sect. 2.3), the SMB inversion is driven by the geometry change between the two DEM epochs. The bedrock is already fixed through Eq. (3); combining it with the observed start-of-window surface DEM gives the thickness that initialises the forward run,

$$H_{t_1}(\mathbf{x}) = z_s^{\text{obs}}(\mathbf{x}, t_1) - z_b(\mathbf{x}). \quad (4)$$

180 By the same construction, combining the fixed bedrock with the observed end-of-window surface DEM returns exactly the DA-calibrated thickness,

$$H^{\text{obs}}(\mathbf{x}, t_2) = z_s^{\text{obs}}(\mathbf{x}, t_2) - z_b(\mathbf{x}) \equiv H_{t_2}(\mathbf{x}), \quad (5)$$

which is therefore the target the forward run is fitted against: the data assimilation's own calibrated state rather than an independent observation.

185 2.4.2 Forward model.

With the derived thickness H_{t_1} (Eq. 4) and the flow scalars (τ_{ref}, A) held fixed, the glacier is evolved from t_1 to t_2 by integrating (1) with an explicit forward-Euler scheme,

$$H^{n+1} = \max(0, H^n + \Delta t^n [smb^n - \nabla \cdot \mathbf{q}^n]), \quad (6)$$

where the non-negativity projection prevents spurious negative thicknesses near the margin. The integration starts from $H^0 =$
190 H_{t_1} and advances until the accumulated time reaches t_2 ; the number of steps N_t is therefore not prescribed but follows from the step sizes selected along the way. Each step is limited by a Courant–Friedrichs–Lewy condition,

$$\Delta t^n = \min\left(c_{\text{cfl}} \Delta x / \max_{\Omega} \|\bar{\mathbf{u}}^n\|, \Delta t_{\text{max}}\right), \quad (7)$$

with $c_{\text{cfl}} = 0.2$ ensuring stability and $\Delta t_{\text{max}} = 1$ yr capping the step where the ice is slow. The flux divergence is evaluated on the staggered grid with a slope-limited scheme, which suppresses spurious oscillations at the steep thickness gradients near the
195 margin. At every step the surface is updated as $z_s^n = z_b + H^n$ and $\bar{\mathbf{u}}$ is recomputed by the emulator from the current (H, z_s) , so the flux divergence evolves with the geometry rather than being held at its initial-snapshot value. Conventional flux-divergence methods (Miles et al., 2021; Kneib et al., 2024), which we retain as a reference in Sect. 3, instead evaluate $smb = \partial H / \partial t + \nabla \cdot \mathbf{q}$ once from the calibrated state.




Figure 3. Mapping the one-dimensional profile to the two-dimensional source term, $smb(\mathbf{x}) = smb(z_s(\mathbf{x}))$, on Argentière. (a) The inferred elevation profile $smb(z)$; (b) the surface elevation $z_s(\mathbf{x})$; and (c) the resulting 2D smb field, obtained by evaluating the profile at the surface elevation of every grid cell.

2.4.3 SMB parametrisation.

200 Rather than treat the smb as an independent unknown at every grid cell, which would leave the inversion severely underdetermined, we parametrise it as a one-dimensional profile $smb(z)$, a function of surface elevation only. This reflects the first-order elevation control on alpine smb and collapses the unknown from $\mathcal{O}(10^4)$ pixels to $N = \mathcal{O}(10)$ elevation bins, which regularises the otherwise severely ill-posed inversion without site-specific smoothing length scales or density assumptions. The profile is discretised on uniformly spaced elevation bins, 200m wide for all three glaciers, by the values $\mathbf{m} = (m_1, \dots, m_N)$, and it is

205 mapped to the two-dimensional source term in (1) by evaluating it at the current surface elevation, $smb(\mathbf{x}, t) = smb(z_s(\mathbf{x}, t))$, with piecewise-linear interpolation between bins (Fig. 3). The parametrisation is deliberately minimal: our aim is an smb that explains the observed geometry change, not the resolution of its fine spatial structure.

2.4.4 Inverse problem.

We seek the smb profile $\mathbf{m}^*$ that minimises the misfit between the simulated end-of-window thickness $H^{\text{sim}}(\cdot, t_2; \mathbf{m})$ and the

210 target $H^{\text{obs}}(\cdot, t_2)$ from (5):

$$\mathcal{J}(\mathbf{m}) = \sqrt{\frac{1}{|\Omega|} \sum_{\mathbf{x} \in \Omega} (H^{\text{sim}} - H^{\text{obs}})^2} + \mathcal{R}(\mathbf{m}), \quad (8)$$

where the first term is the RMS geometry misfit and $\mathcal{R}(\mathbf{m}) = \lambda \sum_i (m_{i-1} - 2m_i + m_{i+1})^2$ is a Tikhonov smoothness penalty on the squared discrete curvature of the profile. The weight λ carries a direct physical reading: it sets how much the smb may change from one elevation bin to the next, and hence how strongly the recovered gradient is smoothed (Fig 4). We use $\lambda = 0.05$




Figure 4. Choice of the smoothness weight λ , shown for Rhône. Each point is one inversion, all sharing the same calibrated dynamical state so that only λ varies. (a) Stronger smoothing lowers the roughness of the profile at the cost of the geometry misfit. (b) Stake RMSE against λ ; the shaded band marks the range over which the recovered profile is effectively unchanged, and the fit degrades only at the extremes. The star is the adopted $\lambda = 0.05$.

for all three glaciers rather than tuning it per site. Varying it over 0.01–0.5 on Rhône moves the geometry misfit by less than 5% while the profile roughness falls by an order of magnitude (Fig. 4a), and the recovered profile is insensitive through the centre of that range: between $\lambda = 0.05$ and 0.1 the equilibrium-line altitude is unchanged and the stake RMSE differs by 0.01 m w.e. yr⁻¹ (Fig. 4b). Only the extremes degrade the fit, the profile tracking noise in the geometry change at $\lambda = 0.01$ and being over-smoothed at $\lambda = 0.5$. The minimisation uses L-BFGS (Liu and Nocedal, 1989) with $M = 10$ correction pairs and a Hager–Zhang line search; the low dimensionality of $\mathbf{m}$ makes the implicit Hessian approximation accurate, and convergence is reached in $\mathcal{O}(10^2)$ function evaluations.

The gradient $\nabla_{\mathbf{m}} \mathcal{J}$ is obtained by reverse-mode automatic differentiation (Baydin et al., 2018) through the entire forward integration (1)–(6): every operation in the forward map—the emulator evaluation, the slope-limited divergence, the explicit Euler update and the piecewise-linear elevation interpolation—is a differentiable TensorFlow (Abadi et al., 2016) operation, so the chain rule applies uniformly across the forward graph and the sensitivity of the end-of-window geometry to each bin value is propagated exactly through every time step.

The principal cost is memory. Automatic differentiation must retain intermediate quantities from the forward run in order to traverse it backwards, and a decade-long integration exceeds the memory of a single GPU. We therefore use gradient checkpointing (Griewank and Walther, 2000; Chen et al., 2016): instead of storing the internal quantities of every step, we keep only the glacier state (H, z_s) at the start of each step and recompute that step’s internals when the backward pass reaches it. The saving is substantial because the emulator evaluation—by far the heaviest part of a step—is never stored, only replayed; the price is one additional evaluation of each step per gradient.




Figure 5. Calibrated data-assimilation state on Argentière at the calibration epoch t_2 , for the adopted velocity-only run ($\tau_{\text{ref}} = 0.199 \text{ MPa}$). (a) Calibrated ice thickness H_{t_2} ; (b) observed surface speed u_s^{obs} (Millan et al., 2022); (c) modelled surface speed u_s^{mod} , in close agreement with the observations.

3 Results

We evaluate the method on three Alpine glaciers of contrasting size—Argentière, Great Aletsch and Rhône—each processed through the same glacier-agnostic two-step pipeline: the data assimilation inverts the ice thickness while the Arrhenius factor A and the basal-motion parameter τ_{ref} are fixed by an ensemble search to the scalar that minimises the surface-velocity misfit, after which the smb is inferred. No per-glacier retuning of the regularisation, optimiser or parametrisation is applied, and none of the three uses a thickness survey. Argentière additionally carries the state-of-the-art distributed reconstruction of Kneib et al. (2024), against which we compare our velocity-only result directly.

3.1 Calibrated ice-dynamics states

For every glacier the data assimilation converges within $\mathcal{O}(10^2)$ L-BFGS iterations to a thickness field whose modelled surface velocity matches the observations to a root-mean-square residual of order $1\text{--}10 \text{ m yr}^{-1}$ —1.0 on Argentière, 10.3 on Great Aletsch and 7.8 on Rhône at the adopted (τ_{ref}, A)—one to two orders of magnitude below typical trunk velocities (locally exceeding $100\text{--}150 \text{ m yr}^{-1}$). The calibrated state fixes the bedrock through Eq. (3), from which the start-epoch geometry H_{t_1} is derived to initialise the transient forward run (Sect. 2.4).

Figure 5 shows the calibrated state on Argentière: the recovered ice-thickness field (a), and the close match between the observed (b) and modelled (c) surface speed, whose differences are small and spatially unstructured.

As a final check on the calibrated state itself, Appendix A2 integrates both candidate starting states on Argentière (H_{t_1}, H_{t_2}) forward for one year under zero smb : both drift with a near-zero mean (-0.3 and -0.1 m), so neither is spontaneously gaining or losing mass under the calibrated dynamics—confirming the fixed-substrate assumption of Sect. 2.1 is not masking a systematic bias that the smb inversion would otherwise have to absorb.

Table 2. Validation of the inferred SMB against the GLACIOCLIM stake network on Argentière Glacier (main trunk + Améthystes tributary) over 2012–2021. The three baselines are the modelling approaches of Kneib et al. (2024), evaluated against the same stakes; all three use ice-thickness data, whereas our result uses surface velocity only.

Method	RMSE [m w.e. yr ⁻¹]
F2019 (Kneib et al., 2024)	0.50
SIA (Kneib et al., 2024)	0.67
IGM-distributed (Kneib et al., 2024)	0.85
Ours (velocity only, scalar τ_{ref})	0.40

3.2 Inferred SMB and validation

Holding the calibrated state (the bedrock fixed via H_{t_2} , and τ_{ref} , A) and the derived start-epoch thickness H_{t_1} fixed, we run the transient forward simulation over each glacier’s window and invert for the elevation-resolved *smb* profile $smb(z)$ that minimises the geometry-misfit cost (Eq. 8). On all three glaciers the retrieved profile increases monotonically across the full altitudinal range, from strongly negative values at the tongue to a positive accumulation regime in the upper basin.

On Argentière (2012–2021) the profile reproduces the observed altitudinal pattern of the GLACIOCLIM stakes with a root-mean-square error of 0.40 m w.e. yr⁻¹ (Fig. 6a, Table 2). This is obtained with the generic velocity-only pipeline—no thickness survey, no per-glacier tuning, no *post hoc* smoothing—yet it surpasses the best physically based reconstructions of Kneib et al. (2024) on the same stakes, whose flux-gate F2019 estimate gives 0.50 and whose SIA and IGM-distributed approaches give 0.67 and 0.85 m w.e. yr⁻¹, all of which additionally require ice-thickness data. Because the inversion variable is one-dimensional, horizontal heterogeneity within an elevation band—notably the avalanche-fed accumulation hotspots reported by Kneib et al. (2024)—is not resolved; this deliberate trade-off, which buys a tuning-free and transferable inversion, is taken up in Sect. 4.

The same pipeline, applied without modification to Great Aletsch (2009–2017) and Rhône (2007–2016), yields equally smooth, monotonic profiles that reproduce the independent stake averages—two continuous-record locations on Aletsch and five on Rhône (Sect. 2.2)—with root-mean-square errors of 0.52 and 0.60 m w.e. yr⁻¹ (Fig. 6b,c). This fit without any thickness constraint reads less as a coincidence than as evidence that the velocity-calibrated dynamical state carries the elevation structure of the SMB; whether the robustness holds across a larger and more varied sample than three well-monitored glaciers is an open question we return to in Sect. 4.

Figure 6 makes concrete the mass-conservation argument of Sect. 1.1. The dashed curve is the SMB the continuity equation yields when the flux divergence is taken as stationary, exactly as in flux-divergence inversions. On every glacier this estimate is markedly noisier than the inferred profile and diverges in the ablation zone, where the stationary divergence is least representative, raising the stake RMSE to 0.98 (Argentière), 1.13 (Aletsch) and 1.53 m w.e. yr⁻¹ (Rhône)—a factor of 2.2 to 2.5 worse than our transient inversion. Because our forward model integrates the mass-conservation law directly and recomputes




Figure 6. Inferred SMB profile (black) for (a) Argentière, (b) Great Aletsch and (c) Rhône, compared against in situ stake measurements and against the continuity estimate $dh/dt + \nabla \cdot \mathbf{q}$ (dashed) obtained with the *stationary* data-assimilation flux divergence. Green diamonds are the continuous-record stake averages ($\pm$ interannual standard deviation), the validation targets (faint grey dots on Aletsch and Rhône are the raw single-year readings). The inferred profile tracks the averages across the whole elevation range (RMSE = 0.40, 0.52 and 0.60 m w.e. yr⁻¹), whereas the stationary-flux estimate departs strongly in the ablation zone (0.98, 1.13 and 1.53).

the flux divergence from the evolving geometry at every step, the inferred profile requires no spatial smoothing and remains mass-consistent throughout.

3.3 Sensitivity to the basal-motion parameter

The basal-motion parameter τ_{ref} is the dynamical quantity least constrained on a glacier without thickness or stake data, where its value rests on the surface-velocity misfit (*velsurf*) alone. It is also intrinsically hard to separate from the ice thickness: a more slippery bed (smaller τ_{ref}) inferred together with thinner ice reproduces almost the same surface velocity—and hence the same ice flux—as a stiffer bed carrying thicker ice, so many (τ_{ref}, H) pairs fit the observations nearly equally well. For the *smb* inversion this ambiguity is not fatal, because what drives the geometry change is not the pair itself but the flux divergence $\nabla \cdot \mathbf{q}$ it produces. To see how the residual freedom in τ_{ref} propagates into the inferred *smb*, we repeated the *smb* inversion across a wide range of τ_{ref} on all three glaciers (Fig. 7), using the independent stake RMSE as ground truth against *velsurf*, the only selector available operationally.

Two features stand out. First, *velsurf* reliably rejects the strongly biased runs: the profiles that depart most from the stakes are also those with the worst velocity misfit (Fig. 7). On each glacier, higher velocity misfit goes with higher stake error (Fig. 8;




Figure 7. Inferred SMB profiles across a range of the basal-motion parameter τ_{ref} for (a) Argentière, (b) Great Aletsch and (c) Rhône; colour encodes τ_{ref} and the adopted profile is drawn bold. Within the low-velocity-misfit basin the profiles coincide in the accumulation area and diverge only at the ablation tongue, so the residual uncertainty in τ_{ref} translates into an equilibrium-line-altitude band (shaded) rather than a whole-profile bias; the runs the velocity misfit rejects (dashed) depart most strongly.

Pearson $r = 0.99$, Spearman $\rho = 0.83$ on Argentière; $r = 0.72$, $\rho = 0.38$ on Aletsch; $r = 0.81$, $\rho = 0.71$ on Rhône). Second,
 290 once the high-misfit runs are removed, the surviving low-*velsurf* runs no longer map one-to-one onto stake error: because a wide range of (τ_{ref}, H) pairs share nearly the same flux, velocity cannot single out one value of τ_{ref} , and on Rhône the stake RMSE even varies non-monotonically among them (0.60 – 1.34 m w.e. yr $^{-1}$). Crucially, this residual freedom is confined to the ablation zone: across the low-misfit runs the accumulation-area profile is essentially unchanged, and only the tongue *smb*—and with it the equilibrium-line altitude—shifts from run to run (Fig. 7).

295 The practical consequence is that the method does not deliver a single “best” profile but a family of low-*velsurf* solutions that converge to the same accumulation-area gradient and fan out only at the tongue (Fig. 7, Fig. 8). For a stake-free glacier this family is the result: the low-misfit runs together define an *smb* and equilibrium-line uncertainty band—narrow in the accumulation area, widening towards the tongue. How this band behaves at regional scale, and how it might be tightened, we take up in Sect. 4.




Figure 8. Data-assimilation surface-velocity misfit (*velsurf* RMSE) against the independent stake RMSE across the τ_{ref} sweep, shown separately for each glacier. Higher velocity misfit is associated with higher stake error on all three (Pearson $r = 0.99, 0.72$ and 0.81 ; Spearman $\rho = 0.83, 0.38$ and 0.71 , for Argentière, Aletsch and Rhône respectively), so the velocity misfit—the only selector available without stakes—carries usable information about the accuracy of the inferred SMB.

300 4 Discussion

Separating surface mass balance from dynamic thinning in an observed surface-elevation change is a notoriously ill-posed problem, yet it is exactly what deriving SMB from remote sensing demands. We have addressed it with a transient, differentiable inversion: a state-of-the-art assimilation of surface ice speed fixes the dynamical state, and a novel automatic-differentiation scheme then inverts the multi-temporal geometry change for the elevation-resolved SMB. On three Alpine glaciers of contrasting size the method reproduces the in situ stake balances to root-mean-square errors of 0.40 (Argentière), 0.52 (Great Aletsch) and 0.60 m w.e. yr⁻¹ (Rhône), surpassing the best physically based reconstruction of Kneib et al. (2024) on the same Argentière stakes with surface velocity alone. Two features set the approach apart from that state of the art. First, the flux divergence is recomputed from the evolving geometry at every time step rather than held at a single stationary snapshot. A stationary estimate evaluates $\nabla \cdot \mathbf{q}$ once, on the geometry of a single epoch, and applies it to an elevation change averaged over a decade—a decade over which the glacier thins, slows and changes its flux. Integrating the mass-conservation equation forward removes that assumption rather than correcting for it, so the inferred SMB is dynamically consistent across the whole window; against the same stakes, the stationary-flux estimate is markedly worse on all three glaciers (Sect. 3).

Second, the inversion needs only surface velocity—no ice-thickness survey and no per-glacier tuning—and proved largely insensitive to the rate factor A , so a single glacier-independent pipeline transfers across all three sites. Beyond validating the inversion, the recovered $smb(z)$ is the quantity mass-balance models are calibrated against: an elevation-resolved balance over a known period is what a temperature-index or energy-balance model must reproduce before it can be driven forward under a climate scenario. Inferring it from geometry alone therefore offers a route to constraining such models—and with them, projections of glacier change and meltwater supply—on the great majority of glaciers that carry no stake record.

The clearest methodological gain is in how mass conservation is enforced. Flux-divergence inversions form the SMB by differentiating the observed thickness–velocity flux and combining it with the elevation-change field; because that divergence is the spatial derivative of a product of two noisy fields, it must be smoothed before use, and the smoothing—as Miles et al. (2021) and Kneib et al. (2024) document—breaks the very mass-conservation constraint the method rests on, leaving a spurious glacier-wide flux-divergence residual of order $0.1\text{--}0.7\text{ m yr}^{-1}$ that they close by redistributing the excess or deficit uniformly across the glacier. Such a homogeneous redistribution is a pragmatic fix rather than a physical statement: it spreads a computational residual over cells that did not produce it. Our formulation removes the need for either step. Because we integrate the mass-conservation equation forward and evaluate the flux divergence on the modelled geometry with a conservative, slope-limited scheme, mass is conserved by construction at every step; the inferred SMB field requires no post-hoc smoothing and no ad-hoc redistribution, and the divergence it implies is guaranteed consistent with the simulated geometry change.

The contrast is visible in our own runs. The flux divergence implied by the calibrated assimilation state is markedly rougher than the thickness field it derives from—grid-scale structure carries three to four times the share of its variance—because differentiating a product of two fields amplifies exactly the wavelengths the assimilation constrains least. Integrating forward damps it: the flux is diffusive in thickness, so by the end of the window the modelled divergence is smooth and spatially coherent where the assimilated one was speckled. The filtering that flux-divergence methods must impose externally is here a consequence of the ice dynamics rather than a step applied afterwards.

The method’s principal limitation is its one-dimensional parametrisation: horizontal variability within an elevation band—such as the avalanche-fed accumulation hotspots resolved by Kneib et al. (2024)—is deliberately not recovered. This buys a transferable, tuning-free inversion and is not intrinsic to the scheme, which admits richer descriptions (an aspect or debris-cover dependence, say) wherever the data warrant them. A second limitation is the thickness–sliding ambiguity of Sect. 3.3: surface velocity constrains the basal-motion parameter only to a band, and the resulting freedom, negligible in the accumulation area, opens into an equilibrium-line spread at the tongue. Two further dependencies are shared with flux-based methods. Accuracy is ultimately set by the input velocity and DEMs: Kneib et al. (2024) find that substituting the coarser global velocity mosaic of Millan et al. (2022) into their reconstruction degrades the stake agreement to $\sim 1.2\text{--}1.3\text{ m w.e. yr}^{-1}$, so the velocity field can itself become limiting—that our transient inversion attains $0.40\text{--}0.60\text{ m w.e. yr}^{-1}$ using precisely that global product is encouraging, but its robustness to velocity error should not be overstated. Finally, like all geodetic estimates the method converts a volume change to a mass balance under an assumed ice density and neglects evolving firn compaction; where the firn column is thick or densifying, this introduces a systematic offset.

Taken together, these properties point towards regional-scale application. Because the pipeline requires no thickness survey, no stake calibration and no site-specific tuning, it can in principle be applied wherever multi-temporal DEMs and a surface-velocity field exist—and both are approaching near-global coverage, through the elevation-change archive of Hugonnet et al. (2021) and expanding velocity products (Millan et al., 2022). Its accuracy will track that of these inputs: substituting that dh/dt product for the dedicated DEM pair on Rhône raises the stake error to $0.77 \text{ m w.e. yr}^{-1}$ (from 0.60), while still recovering the full altitudinal structure and remaining far better than the stationary-flux estimate (Appendix A1). The highest fidelity therefore still comes from dedicated high-resolution DEMs, but the method degrades gracefully rather than failing where only near-global data exist; and the tuning-free character makes a continuous, distributed SMB inversion over many glaciers a realistic next step. Whether the robustness shown here on three well-monitored Alpine glaciers persists across the smaller and more varied population that dominates regional water resources is the key open question, and the natural direction for future work.

5 Conclusions

We have presented a transient, differentiable inversion that recovers the distributed surface mass balance of a glacier from two surface DEMs and a single surface-velocity field, without ice-thickness measurements or per-glacier tuning. Surface velocity is assimilated to fix a dynamically consistent state, and the mass-conservation equation is then integrated forward with automatic differentiation to infer the elevation-resolved SMB whose modelled geometry change best matches observation. Applied through one glacier-independent pipeline to three Alpine glaciers of contrasting size, the method reproduces in situ stake balances to $0.40\text{--}0.60 \text{ m w.e. yr}^{-1}$ —surpassing the state-of-the-art reconstruction of Kneib et al. (2024) on Argentière with surface velocity alone—and a factor of roughly two to two-and-a-half better than the stationary-flux estimates it replaces.

The central advance is that mass conservation is enforced by construction: because the flux divergence is recomputed from the evolving geometry at every step, the inferred SMB needs neither the post-hoc smoothing nor the ad-hoc mass redistribution that flux-divergence methods require, and stays dynamically consistent over the whole observation window. Its accuracy is limited chiefly by the input velocity and DEMs and by the one-dimensional elevation parametrisation, which trades sub-band spatial detail for transferability. Because the method needs no calibration data beyond fields that are approaching near-global coverage, it opens a realistic path towards continuous, regional-scale SMB inversion; establishing its robustness across the smaller and more varied glacier population that dominates regional water resources is the natural next step.

Appendix A: Appendix

A1 Applicability to near-global elevation-change data

The fidelity of the inferred SMB is ultimately set by the quality of the input elevation change (Sect. 1.2). The main results use dedicated two-epoch DEMs, which deliver the highest accuracy but are not available everywhere. Near-global elevation-change products such as that of Hugonnet et al. (2021) would, by contrast, make the method applicable to essentially any glacier. To gauge the associated fidelity penalty in isolation, we repeat the inversion on Rhône—chosen for its denser stake record—

Table A1. Sensitivity of the Rhône inversion to the basal-motion parameter τ_{ref} (at $A = 66.8 \text{ MPa}^{-3} \text{ yr}^{-1}$, the main-text ensemble-search value), comparing the Hugonnet-anchored experiment (Hugonnet-implied 2007 surface) against the dedicated-DEM one (surveyed 2007 surface) at matching τ_{ref} . The two experiments share an identical DA state (2016 anchor, bed, ensemble search) and are validated against the *same* $n = 5$ 2007–2016 stake set, so the two stake-RMSE columns isolate the effect of the 2007 DEM source alone. *velsurf* is identical between the two experiments to the precision shown, since the DA does not depend on the 2007 surface; ELA is from the Hugonnet experiment (the dedicated-DEM values are 60–90 m lower throughout, Fig. A1). Mean improvement: $0.11 \text{ m w.e.yr}^{-1}$, or 9.8% averaged per-row

τ_{ref} (MPa)	velsurf RMSE (m yr ⁻¹)	ELA (Hugonnet) (m a.s.l.)	Stake RMSE (m w.e. yr ⁻¹)	
			Hugonnet	ded. DEMs
0.098	7.19	3041	1.29	1.34
0.132 [†]	7.82	3116	0.77	0.60
0.199	9.45	3061	0.93	0.77
0.744	35.24	3190	2.03	1.91
1.078	37.39	3216	1.93	1.78

keeping *every* other element of the main-text experiment fixed: the data assimilation is anchored on the same real swisstopo 2016 DEM (Sect. 2.3), which fixes the same bed; the window is the same 2007–2016 period; the scalar- τ_{ref} ensemble search is unchanged; and validation uses the *same* $n = 5$ 2007–2016 stake set as the main text. Only the source of the 2007 start-of-window geometry changes: instead of the surveyed swisstopo 2007 DEM, it is derived as $usurf(2016) - \dot{h}_{\text{Hugonnet}} \times 9 \text{ yr}$, using Hugonnet et al. (2021)’s mean 2000–2020 rate. Because the two experiments share an identical DA state, this isolates the effect of the DEM source from that of the validation window. On the ice-covered area, the Hugonnet-implied 2007 surface differs from the surveyed swisstopo DEM by a mean of +3.1 m (RMS 6.0 m).

Using the same velocity-selected coefficient as the main text ($\tau_{\text{ref}} = 0.132 \text{ MPa}$), the Hugonnet-anchored inversion reproduces the stake averages with a root-mean-square error of $0.77 \text{ m w.e. yr}^{-1}$ (Fig. A1)—about 28 % worse than the 0.60 obtained with the dedicated swisstopo DEMs at the same τ_{ref} and identical validation window (Table A1), yet still recovering the full altitudinal structure and remaining well below the stationary-flux continuity estimate ($1.56 \text{ m w.e. yr}^{-1}$, against 1.53 for the dedicated DEMs at the same τ_{ref}). The degradation is concentrated in the accumulation area (Fig. A1b): the equilibrium-line altitude implied by the Hugonnet-anchored profile sits 60–90 m higher than its dedicated-DEM counterpart at every τ_{ref} (Table A1), reflecting the stronger spatial smoothing and larger uncertainty of the near-global product relative to a dedicated DEM pair. This confirms the trade-off anticipated in Sect. 1.2: near-global elevation-change data extend the method to almost any glacier at a fidelity cost bounded by the DEM quality, so dedicated multi-temporal DEMs remain preferable where they exist.




Figure A1. Sliding-coefficient sensitivity of the Rhône inversion for two 2007–2016 experiments sharing an identical DA state (Sect. 2.3, anchored on the real 2016 DEM): (a) the surveyed swisstopo 2007 DEM used in the main text and (b) a 2007 surface implied by applying the Hugonnet et al. (2021) dh/dt product backward from the same 2016 DEM. Each panel shows the inferred SMB profile across a range of τ_{ref} (colour; the profile at the main text’s adopted $\tau_{\text{ref}} = 0.132$ bold with its stake RMSE, and runs with velsurf misfit $> 1.5 \times$ the panel’s best rejected, dashed), the equilibrium-line-altitude band shaded over the retained runs, all validated against the *same* $n = 5$ 2007–2016 stake averages (green diamonds). Both experiments recover the elevation structure, but the Hugonnet-based fan sits higher and fits the stakes less well in the accumulation area, giving a larger stake RMSE at the adopted τ_{ref} (0.77 vs 0.60 m w.e. yr⁻¹).




Figure A2. Mechanical-equilibrium check on Argentière: on-ice thickness drift after integrating one year forward under $smb \equiv 0$, starting from (a) the derived start-epoch state H_{t_1} and (b) the DA-calibrated state H_{t_2} ; (c) the two drift distributions overlaid, with their on-ice RMS.

A2 Mechanical equilibrium of the calibrated dynamical state

395 Sect. 2.1 treats the DA-calibrated bedrock, sliding field and ice-flow emulator as a fixed dynamical substrate and attributes all subsequent thickness change over the inversion window to the free $smb(z)$ profile (Sect. 2.4). That attribution requires the substrate itself to be close to mechanical equilibrium: if it were not, integrating it forward would drift even under zero mass-balance forcing, and that drift would alias into the inferred smb as an initialisation shock (Seroussi et al., 2011). We test this directly on Argentière by integrating each of the two candidate starting states for one year under $smb \equiv 0$ and measuring

400 the resulting on-ice thickness change: the DA-calibrated state H_{t_2} itself (fit directly to the velocity observations, Sect. 2.3), and the derived start-epoch state H_{t_1} (Eq. 4) that the production run actually starts from.

Both states show a near-zero mean drift over the year (-0.10 m for H_{t_2} , -0.33 m for H_{t_1} ; Fig. A2c)—an order of magnitude below the metre-scale annual smb the inversion subsequently recovers—so neither state is spontaneously gaining or losing mass under the calibrated dynamics: there is no systematic bias left for the smb inversion to compensate for. The drift distribution is

405 close to symmetric about zero (Fig. A2c) and concentrated at a handful of margin and tributary-confluence cells (Fig. A2a,b) rather than spread evenly over the glacier, consistent with local relaxation of the discretised stress balance rather than a large-scale imbalance. In absolute terms this local redistribution is not negligible—on-ice RMS = 7.2 m for H_{t_2} and 9.7 m for H_{t_1} (35 % larger), against the 15.8 m RMS thickness change the smb inversion itself fits over the full nine-year window (Sect. 3)—so we report it rather than discount it; but its near-zero mean and localised footprint show it behaves as noise the free $smb(z)$

410 parameters absorb, not as a coherent drift that would bias the inferred mass balance.

References

- Abadi, M., Barham, P., Chen, J., Chen, Z., Davis, A., Dean, J., Devin, M., Ghemawat, S., Irving, G., Isard, M., Kudlur, M., Levenberg, J., Monga, R., Moore, S., Murray, D. G., Steiner, B., Tucker, P., Vasudevan, V., Warden, P., Wicke, M., Yu, Y., and Zheng, X.: TensorFlow: A system for large-scale machine learning, in: 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI 16), pp. 265–283, 2016.
- Akiba, T., Sano, S., Yanase, T., Ohta, T., and Koyama, M.: Optuna: A next-generation hyperparameter optimization framework, in: Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, pp. 2623–2631, <https://doi.org/10.1145/3292500.3330701>, 2019.
- Baydin, A. G., Pearlmutter, B. A., Radul, A. A., and Siskind, J. M.: Automatic differentiation in machine learning: a survey, *Journal of Machine Learning Research*, 18, 1–43, 2018.
- Beraud, L., Cusicanqui, D., Rabatel, A., Brun, F., Vincent, C., and Six, D.: Glacier-wide seasonal and annual geodetic mass balances from Pléiades stereo images: application to the Glacier d’Argentière, French Alps, *Journal of Glaciology*, pp. 1–13, <https://doi.org/10.1017/jog.2022.79>, 2022.
- Bergstra, J., Bardenet, R., Bengio, Y., and Kégl, B.: Algorithms for hyper-parameter optimization, in: Advances in Neural Information Processing Systems 24, pp. 2546–2554, 2011.
- Berthier, E. and Vincent, C.: Relative contribution of surface mass-balance and ice-flux changes to the accelerated thinning of Mer de Glace, French Alps, over 1979–2008, *Journal of Glaciology*, 58, 501–512, <https://doi.org/10.3189/2012JoG11J083>, 2012.
- Bhushan, S., Shean, D., Hu, J.-Y. M., Guillet, G., and Rounce, D. R.: Deriving seasonal and annual surface mass balance for debris-covered glaciers from flow-corrected satellite stereo DEM time series, *Journal of Glaciology*, pp. 1–21, <https://doi.org/10.1017/jog.2024.57>, 2024.
- Bisset, R. R., Dehecq, A., Goldberg, D. N., Huss, M., Bingham, R. G., and Gourmelen, N.: Reversed surface-mass-balance gradients on Himalayan debris-covered glaciers inferred from remote sensing, *Remote Sensing*, 12, 1563, <https://doi.org/10.3390/rs12101563>, 2020.
- Brinkerhoff, D. J., Tober, B. S., Daniel, M., Devaux-Chupin, V., Christoffersen, M. S., Holt, J. W., Larsen, C. F., Fahnestock, M., Loso, M. G., Timm, K. M. F., Mitchell, R. C., and Truffer, M.: The demise of the world’s largest piedmont glacier: a probabilistic forecast, *The Cryosphere*, 19, 2321–2353, <https://doi.org/10.5194/tc-19-2321-2025>, 2025.
- Brun, F., Wagnon, P., Berthier, E., Shea, J. M., Immerzeel, W. W., Kraaijenbrink, P. D. A., Vincent, C., Reverchon, C., Shrestha, D., and Arnaud, Y.: Ice cliff contribution to the tongue-wide ablation of Changri Nup Glacier, Nepal, central Himalaya, *The Cryosphere*, 12, 3439–3457, <https://doi.org/10.5194/tc-12-3439-2018>, 2018.
- Buri, P., Fatichi, S., Shaw, T. E., Miles, E. S., McCarthy, M. J., Fyffe, C. L., Fugger, S., Ren, S., Kneib, M., Jouberton, A., Steiner, J., Fujita, K., and Pellicciotti, F.: Land surface modeling in the Himalayas: on the importance of evaporative fluxes for the water balance of a high-elevation catchment, *Water Resources Research*, 59, <https://doi.org/10.1029/2022WR033841>, 2023.
- Chen, T., Xu, B., Zhang, C., and Guestrin, C.: Training deep nets with sublinear memory cost, arXiv preprint arXiv:1604.06174, 2016.
- Cogley, J. G., Hock, R., Rasmussen, L. A., Arendt, A. A., Bauder, A., Braithwaite, R. J., Jansson, P., Kaser, G., Möller, M., Nicholson, L., and Zemp, M.: Glossary of Glacier Mass Balance and Related Terms, Tech. rep., UNESCO-IHP, Paris, 2011.
- Cuffey, K. M. and Paterson, W. S. B.: *The Physics of Glaciers*, Elsevier, Amsterdam, 4 edn., 2010.
- Dehecq, A., Gourmelen, N., Gardner, A. S., Brun, F., Goldberg, D., Nienow, P. W., Berthier, E., Vincent, C., Wagnon, P., and Trouvé, E.: Twenty-first-century glacier slowdown driven by mass loss in High Mountain Asia, *Nature Geoscience*, 12, 22–27, <https://doi.org/10.1038/s41561-018-0271-9>, 2019.

- Enderlin, E. M., Hamilton, G. S., O'Neel, S., Bartholomaus, T. C., Morlighem, M., and Holt, J. W.: An empirical approach for estimating stress-coupling lengths for marine-terminating glaciers, *Frontiers in Earth Science*, 4, <https://doi.org/10.3389/feart.2016.00104>, 2016.
- 450 GLACIOCLIM: Les GLACIers, un Observatoire du CLIMat, Observatoire des Sciences de l'Univers de Grenoble, <https://glacioclim.osug.fr/>, last access: 11 August 2026, 2026.
- GLAMOS: Swiss Glacier Monitoring Network (GLAMOS), Glacier Monitoring Switzerland, <https://doi.org/10.18750/massbalance.2021.r2021>, 2021.
- GLAMOS: Swiss Glacier Point Mass Balance Observations, release 2023, Glacier Monitoring Switzerland, 455 <https://doi.org/10.18750/massbalance.point.2023.r2023>, 2023.
- Griewank, A. and Walther, A.: Algorithm 799: revolve: an implementation of checkpointing for the reverse or adjoint mode of computational differentiation, *ACM Transactions on Mathematical Software*, 26, 19–45, <https://doi.org/10.1145/347837.347846>, 2000.
- Hansen, P. C.: Analysis of discrete ill-posed problems by means of the L-curve, *SIAM Review*, 34, 561–580, <https://doi.org/10.1137/1034115>, 1992.
- 460 Hartl, L., Schmitt, P., Schuster, L., Helfricht, K., Abermann, J., and Maussion, F.: Recent observations and glacier modeling point towards near-complete glacier loss in western Austria (Ötztal and Stubai mountain range) if 1.5 °C is not met, *The Cryosphere*, 19, 1431–1452, <https://doi.org/10.5194/tc-19-1431-2025>, 2025.
- Herrmann, O., Prasad, V., Zöller, A., Groos, A. R., Cook, S., Sommer, C., and Fürst, J. J.: Brief communication: Inferring glacier equilibrium line altitudes in central Europe with FROST, *EGUsphere* [preprint], <https://doi.org/10.5194/egusphere-2025-5486>, 2025.
- 465 Hugonnet, R., McNabb, R., Berthier, E., Menounos, B., Nuth, C., Girod, L., Farinotti, D., Huss, M., Dussaillant, I., Brun, F., and Käab, A.: Accelerated global glacier mass loss in the early twenty-first century, *Nature*, 592, 726–731, <https://doi.org/10.1038/s41586-021-03436-z>, 2021.
- Jóhannesson, T., Raymond, C., and Waddington, E.: Time-scale for adjustment of glaciers to changes in mass balance, *Journal of Glaciology*, 35, 355–369, <https://doi.org/10.3189/S002214300000928X>, 1989.
- 470 Juvet, G.: Inversion of a Stokes glacier flow model emulated by deep learning, *Journal of Glaciology*, 69, 13–26, <https://doi.org/10.1017/jog.2022.41>, 2023.
- Juvet, G. and Cordonnier, G.: Ice-flow model emulator based on physics-informed deep learning, *Journal of Glaciology*, 69, 1941–1955, <https://doi.org/10.1017/jog.2023.73>, 2023.
- Juvet, G., Finley, B., Gregov, T., and Rosier, S.: IGM: an accessible, modular, differentiable, and GPU-accelerated high-order ice flow 475 model, *EarthArXiv*, <https://doi.org/10.31223/X5GB6G>, preprint, 2026.
- Kamb, B. and Echelmeyer, K. A.: Stress-gradient coupling in glacier flow: I. Longitudinal averaging of the influence of ice thickness and surface slope, *Journal of Glaciology*, 32, 267–284, <https://doi.org/10.3189/S0022143000015604>, 1986.
- Kneib, M., Miles, E. S., Buri, P., Fugger, S., McCarthy, M., Shaw, T. E., Chuanxi, Z., Truffer, M., Westoby, M. J., Yang, W., and Pellicciotti, F.: Sub-seasonal variability of supraglacial ice cliff melt rates and associated processes from time-lapse photogrammetry, *The Cryosphere*, 480 16, 4701–4725, <https://doi.org/10.5194/tc-16-4701-2022>, 2022.
- Kneib, M., Dehecq, A., Gilbert, A., Basset, A., Miles, E. S., Juvet, G., Jourdain, B., Ducasse, E., Beraud, L., Rabatel, A., et al.: Distributed surface mass balance of an avalanche-fed glacier, *EGUsphere*, 2024, 1–35, <https://doi.org/10.5194/egusphere-2024-1733>, 2024.
- Liu, D. C. and Nocedal, J.: On the limited memory BFGS method for large scale optimization, *Mathematical Programming*, 45, 503–528, <https://doi.org/10.1007/BF01589116>, 1989.

- 485 Marzeion, B., Jarosch, A. H., and Hofer, M.: Past and future sea-level change from the surface mass balance of glaciers, *The Cryosphere*, 6, 1295–1322, <https://doi.org/10.5194/tc-6-1295-2012>, 2012.
- Miles, E., McCarthy, M., Dehecq, A., Kneib, M., Fugger, S., and Pellicciotti, F.: Health and sustainability of glaciers in High Mountain Asia, *Nature Communications*, 12, 2868, <https://doi.org/10.1038/s41467-021-23073-4>, 2021.
- Millan, R., Mouginit, J., Rabatel, A., and Morlighem, M.: Ice velocity and thickness of the world’s glaciers, *Nature Geoscience*, 15, 124–129, <https://doi.org/10.1038/s41561-021-00885-z>, 2022.
- 490 Peltó, B. M. and Menounos, B.: Surface mass-balance gradients from elevation and ice flux data in the Columbia Basin, Canada, *Frontiers in Earth Science*, 9, 675 681, <https://doi.org/10.3389/feart.2021.675681>, 2021.
- Radić, V., Bliss, A., Beedlow, A. C., Hock, R., Miles, E., and Cogley, J. G.: Regional and global projections of twenty-first century glacier mass changes in response to climate scenarios from global climate models, *Climate Dynamics*, 42, 37–58, [https://doi.org/10.1007/s00382-](https://doi.org/10.1007/s00382-013-1719-7)
- 495 013-1719-7, 2014.
- Rosier, S. and Juvet, G.: A deep-learning ice-flow emulator for transient glacier modelling, *Journal of Glaciology*, preprint, 2026.
- Rounce, D. R., Hock, R., and Shean, D. E.: Glacier mass change in High Mountain Asia through 2100 using the open-source Python glacier evolution model (PyGEM), *Frontiers in Earth Science*, 7, 331, <https://doi.org/10.3389/feart.2019.00331>, 2020a.
- Rounce, D. R., Khurana, T., Short, M. B., Hock, R., Shean, D. E., and Brinkerhoff, D. J.: Quantifying parameter uncertainty in a large-
500 scale glacier evolution model using Bayesian inference: application to High Mountain Asia, *Journal of Glaciology*, 66, 175–187, <https://doi.org/10.1017/jog.2019.91>, 2020b.
- Schmitt, P., Maussion, F., Goldberg, D. N., and Gregor, P.: AGILE v0.1: the open global glacier data assimilation framework, *Geoscientific Model Development*, 19, 1301–1319, <https://doi.org/10.5194/gmd-19-1301-2026>, 2026.
- Schuster, L., Rounce, D. R., and Maussion, F.: Glacier projections sensitivity to temperature-index model choices and calibration strategies, *Annals of Glaciology*, 64, 293–308, <https://doi.org/10.1017/aog.2023.57>, 2023.
- 505 Seroussi, H., Morlighem, M., Rignot, E., Larour, E., Aubry, D., Ben Dhia, H., and Kristensen, S. S.: Ice flux divergence anomalies on 79north Glacier, Greenland, *Geophysical Research Letters*, 38, L09 501, <https://doi.org/10.1029/2011GL047338>, 2011.
- swisstopo: Federal Office of Topography swisstopo, <https://www.swisstopo.admin.ch>, last access: 11 August 2026, 2026.
- Van Tricht, L., Huybrechts, P., Van Breedam, J., Vanhulle, A., Van Oost, K., and Zekollari, H.: Estimating surface mass balance patterns from
510 unoccupied aerial vehicle measurements in the ablation area of the Morteratsch–Pers glacier complex (Switzerland), *The Cryosphere*, 15, 4445–4464, <https://doi.org/10.5194/tc-15-4445-2021>, 2021.
- Weertman, J.: On the sliding of glaciers, *Journal of Glaciology*, 3, 33–38, <https://doi.org/10.3189/S0022143000024709>, 1957.
- WGMS: Fluctuations of Glaciers (FoG) Database, World Glacier Monitoring Service, Zurich, Switzerland, <https://doi.org/10.5904/wgms-fog-2025-02b>, 2025.
- 515 Zekollari, H. and Huybrechts, P.: On the climate–geometry imbalance, response time and volume–area scaling of an alpine glacier: insights from a 3-D flow model applied to Vadret da Morteratsch, Switzerland, *Annals of Glaciology*, 56, 51–62, <https://doi.org/10.3189/2015AoG70A921>, 2015.
- Zeller, L., McGrath, D., Sass, L., O’Neel, S., McNeil, C., and Baker, E.: Beyond glacier-wide mass balances: parsing seasonal elevation change into spatially resolved patterns of accumulation and ablation at Wolverine Glacier, Alaska, *Journal of Glaciology*, 69, 87–102, <https://doi.org/10.1017/jog.2022.46>, 2023.
- 520

Zemp, M., Huss, M., Thibert, E., Eckert, N., McNabb, R., Huber, J., Barandun, M., Machguth, H., Nussbaumer, S. U., Gärtner-Roer, I., Thomson, L., Paul, F., Maussion, F., Kutuzov, S., and Cogley, J. G.: Global glacier mass changes and their contributions to sea-level rise from 1961 to 2016, *Nature*, 568, 382–386, <https://doi.org/10.1038/s41586-019-1071-0>, 2019.